

REQUEST FOR PROPOSAL (RFP)

Real-Time Enterprise Vector Search & Autonomous LLM Infrastructure

Billion-scale embedding retrieval, private VPC fine-tuning, and multi-tenant agent execution

ISSUING ORGANIZATION & KEY DETAILS

Issuing Entity:	Beacon Data Intelligence (beacondata.ai)
Primary Contact:	Liam O'Connor, Chief Data Officer
Direct Inquiries:	liam.oconnor@beacondata.ai
Publication Date:	October 2026
Estimated Contract Value:	\$3,800,000 USD
Procurement Protocol:	Competitive Sealed Proposal (Phase-Gate Evaluation)

NOTICE: The information contained in this Request for Proposal is confidential and proprietary to Beacon Data Intelligence. By accepting this document, prospective vendors agree to hold all technical architectures, commercial terms, and proprietary data in strict confidence under non-disclosure provisions.

1. Project Summary

Beacon Data Intelligence powers enterprise semantic discovery and generative AI knowledge engines for Fortune 500 corporations. We are expanding our core infrastructure to support billion-scale vector indices and autonomous multi-agent reasoning pipelines.

2. Technical & Scaling Requirements

- Billion-Scale Vector Indexing: Distributed HNSW / ScaNN vector indexing supporting 2+ billion high-dimensional embeddings with <20ms p99 query latency.
- Private VPC Deployment: Complete customer isolation with zero data egress to public model APIs.
- Dynamic Re-ranking: Hybrid sparse/dense retrieval combining BM25 and neural cross-encoders.
- Agentic Orchestration: Low-latency framework for hosting autonomous reasoning agents with memory caching and tool-use verification.

3. Governance & Performance Guarantees

- SOC 2 Type II, ISO 27001, and ISO 42001 (Artificial Intelligence Management System) compliance.
- Complete role-based access control (RBAC) ensuring documents retain access permissions down to paragraph chunks.

4. Timeline & Proposal Submission

- Submit technical proposals and benchmark proof to liam.oconnor@beacondata.ai by November 28, 2026.
- Benchmarking bake-off scheduled for early December 2026.